

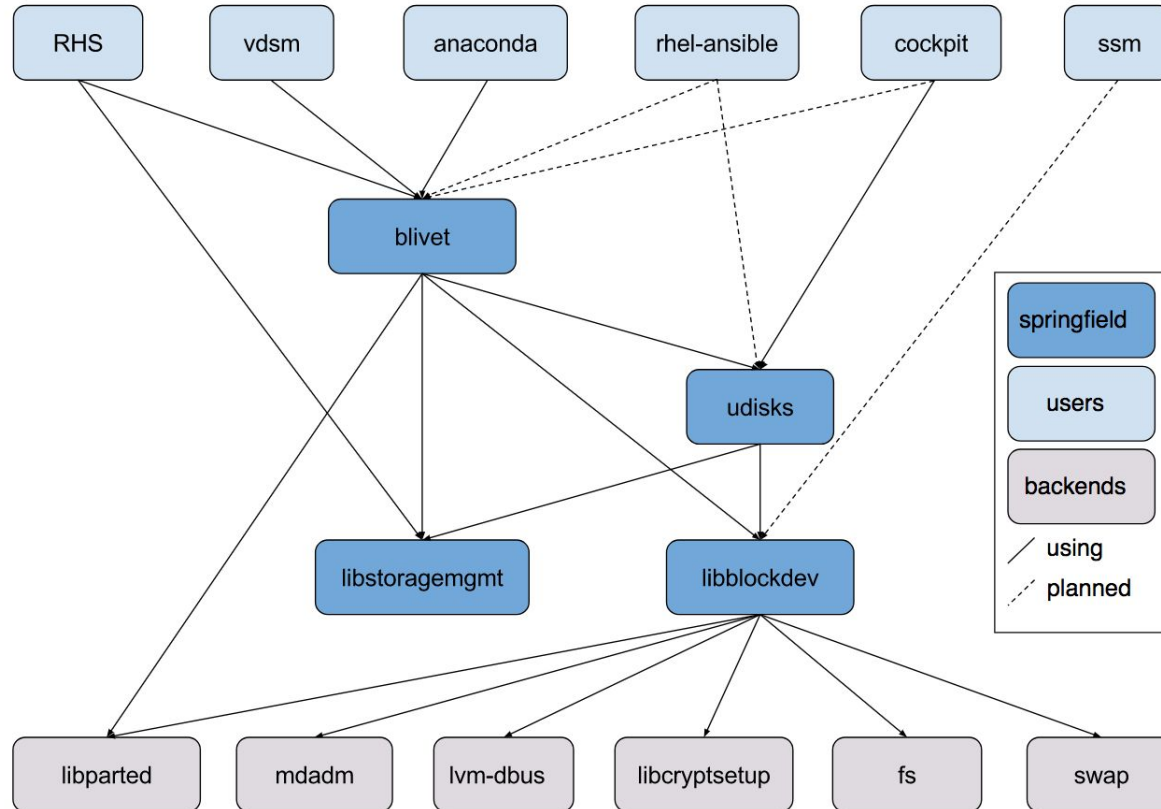
Project Springfield

A smart and automated storage

What is Springfield?

- A collaborative effort to solve certain problems across projects.
 - No “springfield.so” or any binary/executable, but a joint kernel/userland project.
 - “Officially” right now: [udisks](#), [libblockdev](#), [blivet](#), [libstoragemgmt](#)
 - Other projects cooperating more loosely or in early development/planning stages.
- Some issues we are trying to solve:
 - Good for human \neq good for automation (the UNIX way unfit for all?)
 - Informations are fragmented over multiple sources:
 - Redundant queries.
 - Higher CPU usage.
 - No good notification system for storage.
 - Complex APIs, not easy to use by application developers.

Project interactions



What the specific projects do?

- **Blivet**
 - Python module for storage configuration.
 - Can model arbitrary number of changes in memory before committing them.
- **udisks**
 - Daemon providing DBus API for storage configuration and monitoring.
- **libblockdev**
 - A library wrapping around various fs-utils and storage tools and providing an unified API.
 - Low-level.
- **libstoragemgmt**
 - Manages a controller-based storage (e.g. NFS).

Discussion time

Three main topics:

- Filesystem API
- Mount API
- Storage Information Provider

Filesystem API - libblockdev

- I asked about the FS API few weeks ago at linux-fsdevel list.
 - Some interest from potential users.
 - Maintainers: Chicken/egg issue.
 - So we reached for something already there that is “good enough”.
- Works right now (C library with Python bindings).
- Does the screenscraping for you.
- All-in-one, maybe too much dependencies (RAID, ...).
- Is this enough, is it useful for anyone else?
 - E.g. not everyone wants to depend on a 3rd party library, but we probably can't do anything else until there is enough of users of this...

Mount API 1/2

- A single source, data right from kernel and supporting large mount tables.
 - A kernel module + userspace library
- Why not just use libmount as it is?
 - A terrible performance with thousands/ten of thousands volumes (containers, autofs, etc.)
 - For example, starting autofs with a direct mount map of 15,000 entries several key processes consume all available CPU (to name a few of these, systemd, udisksd, gvfs-udisks2-volume-monitor and gvfs-trash). And any change does it again.
 - API is too low-level (ideally we want ~5 lines of code, excluding option setup.)

Mount API 2/2

- Mount namespaces complicates everything.
 - /proc is currently the only way to get namespace-specific information.
 - Maintaining a database of namespaces and mounts is too costly even for moderate load.
- So we can't have a daemon, only a library, right?
- Getting incremental changes of mount table straight from the kernel is important for performance, but it seems like a really difficult thing to do.
 - How to split it? Timestamp and co. is probably too costly/complicated.
 - So a notification to all subscribed listeners on every change?

Storage Information Provider 1/2

- Motivation: no reliable way to detect things like “FS is full”
 - Stacked layers, thin provisioning, deduplication... → user gets write failures for no apparent reason.
 - Notifications, better structured errors able to explain what and where happened.
- Really early preview: <https://github.com/raven-au/usip>
 - Really, really early

Storage Information Provider 2/2

- How to calculate a relative path to a provided root struct? (Namespaces again...)
 - Some functions in kernel can do it, but they are not exported.
 - It probably needs to be a standalone module - it doesn't fit into any fs code or VFS.
 - Workqueue tasks are unconditionally created with a shared current->fs so set_fs_root() can't really be done in a worker, so we need to use kernel_thread(), but is not exported too. Would exporting any of these lead to abuse of mount name spaces? (Probably yes... or other issues.)
 - What are the implications of cloning a workqueue worker thread without setting CLONE_FS?

Thank you

Do you know about a project that would benefit from the cooperation? Tell us!

Contact:

- <https://springfield-project.github.io/>
- springfield@sourceware.org

These slides:

- <https://springfield-project.github.io/lfs2018/>

Me:

- Jan Ťulák - jtulak@redhat.com